

Most Active Stocks
(<https://www.fool.com/market/most-active-stocks/>)

Daily Stock Gainers
(<https://www.fool.com/market/daily-stock-gainers/>)

Daily Stock Losers
(<https://www.fool.com/market/daily-stock-losers/>)

EARNINGS CALL TRANSCRIPT

By Motley Fool
Transcribing
(/author/20032/)
– Nov 19, 2025
at 6:18PM EST

✧ Summarize with AI ✓

Advertisement

(https://ads.freestar.com/?utm_campaign=branding&utm_medium=display&utm_source=fool.com&utm_content=foolcom_leaderboard_atf_article)

Wednesday, November 19, 2025 at 5 p.m. ET

The Motley Fool.

Transcripts

▲ S&P 500 +189% | ▲ Stock Advisor +938% JOIN THE MOTLEY FOO

- President and Chief Executive Officer Jensen Huang 527423-a535-4519-a07f-

20014582e03e&impression=24494911-cf57-4c0c-b586-image source: The Motley

Executive Vice President and Chief Financial Officer —

ACCESSIBILITY
(#)

LOG
IN

65% off* St Patty's Day! › ([https://api.](https://api.HELP)
apikey=35527428-a535-4519-a07f-

Colette Kress (Our Services ▾ Stock Market News ▾ How to Invest ▾ Retirement ▾ Personal Finance ▾ About Us ▾)

- Vice President, Investor Relations – Toshiya Hari

Need a quote from a Motley Fool analyst? Email pr@fool.com

Risks

- **China revenue constraints** – Colette Kress said, "Sizable purchase orders never materialized in the quarter due to geopolitical issues and the increasingly competitive market in China," directly impacting data center compute product shipments to China.
- **Rising input costs** – Colette Kress stated, "input costs are on the rise but we are working to hold gross margins in the mid-seventies," highlighting margin pressure into fiscal 2027.
- Inventory grew 32% quarter over quarter while supply commitments increased 63% sequentially.

Takeaways

- **Total revenue** – \$57 billion, up 62% year over year, with sequential growth of \$10 billion or 22%.
- **Data center revenue** – Record \$51 billion, representing a 66% year-over-year increase, driven mainly by strong AI infrastructure demand.
- **Networking revenue** – \$8.2 billion, up 162% year over year, with growth in NVLink, InfiniBand, and Spectrum X Ethernet.
- **Compute segment** – Grew 56% year over year in Q3 (ended Oct. 26, 2025), driven principally by the ramp of GB 300 GPUs, while Hopper platform delivered approximately \$2 billion in segment revenue in Q3.
- **Gaming revenue** – \$4.3 billion, up 30% year over year, attributed to heightened gamer demand.
- **Professional visualization revenue** – \$760 million, up 56% from the prior year, with DGX Spark cited as a key contributor.
- **Automotive revenue** – \$592 million, up 32% year over year, primarily driven by self-driving solutions.
- **GAAP gross margin** – 73.4%, while non-GAAP gross margin reached 73.6%, both exceeding previous outlook due to data center mix, improved cycle time, and cost structure.
- **Operating expenses** – GAAP operating expenses rose 8% sequentially; non-GAAP

STOCKS MENTIONED

3743-4102-8a3c-04292cbd2fcd&url=https%3A%2F%2Fsa-65-off-spd-11243%26source%3Di%3Dfoolcom-flash-sales%26ftm_veh%3Dfoolcom_nav_b

Nvidia
NASDAQ:
NVDA

Motley Fool Stock Advisor's Latest Pick
Get Access (https://api.fool.com/infotron/infotra4519-a07f-20014582e03e&impression=f9d5c9a997feeb0d214&url=https%3A%2F%2Fwww.fool.com/v-sp-wdg-art%3Fsource%3Disafcmwdg0000003%26ftm_veh%3Darticle_sidebar_small%3Dfoolcom-nav_b)

**Average returns of all recommendations since inception. Cost basis and return based on previous market day close.*

Ad

CHASE

Enjoy \$2,700 in annual value with Chase Sapphire Reserve

Get annual travel credits, airport lounge access, and 3x points on travel and dining - plus 125,000 bonus points after you spend \$6,000 in purchases in the first 3 months. Terms apply. Member FDIC.

Apply Now >

On Chase's Secure Website.
[*Rates & Fees](#)

operating expenses increased 11% sequentially, mainly due to infrastructure compute, compensation, and engineering/development costs.

- **Inventory and supply commitments** -- Inventory increased 32% and supply commitments rose 63% quarter over quarter, as the company prepared for growth.
- **Q4 revenue outlook** -- Expected total revenue of \$65 billion plus or minus 2%, implying 14% sequential growth at the midpoint, driven by continued Blackwell architecture momentum.
- **Gross margin guidance** -- Expected mid-seventies gross margins for both GAAP and non-GAAP, despite rising input costs.
- **AI platform demand** -- Management highlighted fully utilized GPU installed base and "clouds are sold out," indicating persistent supply-demand imbalance.
- **Blackwell platform** -- GB 300 made up roughly two-thirds of total Blackwell revenue; GB 300 is now leading the product transition with broad customer adoption.
- **China data center assumptions** -- Management stated, "we are not assuming any data center compute revenue from China" in the Q4 outlook.
- **Strategic partnerships and investments** -- Announced deals with partners including AWS, Humane, Suzuki, **Intel** (NASDAQ: INTC), **Arm** (NASDAQ: ARM), and Anthropic; up to 5 million GPUs associated with new AI factory projects.
- **Rubin platform** -- On schedule for 2026 ramp, with first silicon delivered and a focus on backward compatibility and rapid ecosystem adoption.
- **Performance leadership** -- Management cited 5x faster time to train vs. Hopper using Blackwell Ultra and 10x higher performance per watt and 10x lower cost per token versus H200 on DeepSeek r1 benchmarks.
- **Strategic investments** -- Continued investment in AI model builders such as OpenAI and Anthropic to deepen ecosystem reach and performance optimization.
- **Supply chain expansion** -- First Blackwell wafer produced on U.S. soil in partnership with **TSMC** (NYSE: TSM); ongoing efforts to broaden supply redundancy and resilience.

Summary

Nvidia (NVDA (/quote/nasdaq/nvda/) 0.35%) reported revenue of \$57 billion with significant growth across all business segments, particularly in data center operations. Management highlighted visibility to a half a trillion dollars in Blackwell and Rubin revenue from the start of this year through the end of calendar year 2026. The GPU installed base was described as fully utilized and "the clouds are sold out," reinforcing

Advertisement
SPONSORED CONTENT **Economic** Advertise Here ▶



Which State Taxes Matter to You?

Charles Schwab



What Is Day Trading? Rules, Risks, and More

Charles Schwab




Natural-born traders, TradeStation was built for you. Start trading.

TradeStation



With E*TRADE, seamlessly trade across web, desktop, and mobile

E*TRADE from Morgan Stanley



8 Ways to Prepare for Tax Filing Season

Charles Schwab

exceptionally strong demand conditions. Networking revenue rose 162% year over year, and GB 300 shipments have overtaken prior Blackwell products, indicating rapid customer adoption of next-generation architectures. Despite geopolitical constraints limiting shipments to China, the company expects gross margin stability in the mid-seventies for the coming year, even as input costs rise and inventory levels increase. Anticipated Q4 revenue of \$65 billion, up 14% sequentially at the midpoint, reflects ongoing momentum in AI infrastructure buildout fueled by landmark deals and broad-based enterprise adoption.

- Colette Kress said, "we shipped \$50 billion this quarter," positioning the company toward a \$500 billion opportunity for Blackwell and Rubin platforms through 2026, with aggregate future demand likely to increase.
- Jensen Huang described three fundamental platform shifts—CPU to GPU acceleration, generative AI mainstreaming, and agentic AI emergence—as core drivers of multi-year infrastructure investment.
- Company collaboration highlights included a strategic partnership with Anthropic, driving adoption of Nvidia architecture for the first time and targeting one gigawatt of compute capacity, scaling upwards.
- Customers such as AWS and Humane announced plans to deploy up to 150,000 AI accelerators, with xAI and Humane co-developing a flagship 500 megawatt data center, illustrating tangible high-scale deployments.
- Management cited broad industry engagement, with new platform launches and deeper integrations across hyperscalers, enterprise software providers, and robotics innovators propelling the CUDA ecosystem's expansion.

Industry glossary

- **CUDA:** Nvidia's parallel computing platform and programming model enabling GPUs to accelerate specialized applications for artificial intelligence and scientific computing.
- **Blackwell:** Nvidia's current-generation AI GPU architecture, including products such as GB 200 and GB 300, optimized for advanced AI workloads.
- **Rubin:** The company's forthcoming data center AI platform, positioned as the next major architecture after Blackwell with backward compatibility and enhanced performance.
- **NVLink:** Nvidia's proprietary high-speed interconnect enabling fast data transfer across GPUs within a computing cluster.
- **Spectrum X Ethernet:** High-performance Ethernet switching family from Nvidia, tailored for AI and accelerated data center networking.

- **Agentic AI:** Artificial intelligence systems capable of autonomous reasoning, planning, and tool use, as described by management as a new wave in computing.
- **MLPerf:** Industry-standard AI benchmarking suite cited for performance measurement of training and inference in machine learning models.

Full Conference Call Transcript

Toshiya Hari: Good afternoon, everyone, and welcome to NVIDIA Corporation's conference call for the 2026. With me today from NVIDIA Corporation are Jensen Huang, president and chief executive officer, and Colette Kress, executive vice president and chief financial officer. I'd like to remind you that our call is being webcast live on NVIDIA Corporation's 2026. The content of today's call is NVIDIA Corporation's property. It cannot be reproduced or transcribed without our prior written consent. During this call, we may make forward-looking statements based on current expectations. These are subject to a number of significant risks and uncertainties, and our actual results may differ materially.

For a discussion of factors that could affect our future financial results and business, please refer to the disclosure in today's earnings release, our most recent forms 10-K and 10-Q, and the reports that we may file on Form 8-K with the Securities and Exchange Commission. All our statements are made as of today, 11/19/2025, based on information currently available to us. Except as required by law, we assume no obligation to update any such statements. During this call, we will discuss non-GAAP financial measures. You can find a reconciliation of these non-GAAP financial measures to GAAP financial measures in our CFO commentary, which is posted on our website. With that, let me turn the call over to Colette.

Advertisement

Colette Kress: Thank you, Toshiya. We delivered another outstanding quarter with revenue of \$57 billion, up 62% year over year and a record sequential revenue growth of \$10 billion or 22%. Our customers continue to lean into three platform shifts fueling exponential growth for accelerated computing, powerful AI models, and agentic applications. Yet we are still in the early innings of these transitions that will impact our work across every industry. Currently, we have visibility to a half a trillion dollars in Blackwell and Rubin revenue from the start of this year through the end of calendar year 2026.

By executing our annual product cadence and extending our performance leadership through full stack design, we believe NVIDIA Corporation will be the superior choice for the \$3 to \$4 trillion in annual AI infrastructure build we estimate by the end of the decade. Demand for AI infrastructure continues to exceed our expectations. The clouds are sold out, and our GPU installed base, both new and previous generations,

including Blackwell, Hopper, and Ampere, is fully utilized. Record Q3 data center revenue of \$51 billion increased 66% year over year, a significant feat at our scale.

Compute grew 56% year over year driven primarily by the GB 300 ramp while networking more than doubled given the onset of NVLink scale up and robust double-digit growth across Spectrum X Ethernet and Quantum X InfiniBand. The world hyperscalers, a trillion-dollar industry, are transforming search recommendations, and content understanding from classical machine learning to generative AI. NVIDIA CUDA excels at both and is the ideal platform for this transition, driving infrastructure investment measured in hundreds of billions of dollars. At Meta, AI recommendation systems are delivering higher quality and more relevant content, leading to more time spent on apps such as Facebook and Threads.

Analyst expectations for the top CSPs and hyperscalers in 2026 aggregate CapEx have continued to increase and now sit roughly at \$600 billion, more than \$200 billion higher relative to the start of the year. We see the transition to accelerated computing and generative AI across current hyperscale workloads contributing toward roughly half of our long-term opportunity. Another growth pillar is the ongoing increase in compute spend driven by foundation model builders such as Anthropic, Mistral, OpenAI, Reflection, Safe Superintelligence, Thinking Machines Lab, and xAI. All scaling, compute aggressively to scale intelligence. The three scaling laws pretraining, post-training, and inference remain intact.

In fact, we see a positive virtuous cycle emerging whereby the three scaling laws and access to compute are generating better intelligence and in turn increasing adoption and profits. OpenAI recently shared that their weekly user base has grown to 800 million. Enterprise customers have increased to 1 million, and their gross margins were healthy. Well, Anthropic recently reported that its annualized run rate revenue has reached \$7 billion as of last month, up from \$1 billion at the start of the year. We are also witnessing a proliferation of agentic AI across various industries and tasks.

Advertisement

Companies such as Cursor Anthropic, Open Evidence, Epic, and Abridge are experiencing a surge in user growth as they supercharge the existing workforce, delivering unquestionable ROI for coders and healthcare professionals. The world's most important enterprise software platforms like ServiceNow, CrowdStrike, and SAP are integrating NVIDIA Corporation's accelerated computing and AI stack. Our new partner, Palantir, is supercharging the incredibly popular ontology platform with NVIDIA CUDA X libraries and AI models for the first time. Previously, like most enterprise software platforms, Ontology runs only on CPUs. Lowe's is leveraging the platform to build supply chain agility, reducing costs, and improving customer satisfaction. Enterprises broadly are leveraging AI to boost productivity, increase efficiency, and reduce cost.

RBC is leveraging agentic AI to drive significant analysts' productivity, slashing report generation time from hours to minutes. AI and digital twins are helping Unilever accelerate content creation by 2x and cut costs by 550%. And Salesforce's engineering team has seen at least 30% productivity increase in new code development after adopting Cursor. This past quarter, we announced AI factory and infrastructure projects amounting to an aggregate of 5 million GPUs. This demand spans every market CSPs, sovereigns, modern builders, enterprises, and supercomputing centers includes multiple landmark build outs. XAI's Colossus two, the world's first gigawatt scale data center. Lilly's AI factory for drug discovery, the pharmaceutical industry's most powerful data center.

And just today, AWS and Humane expanded their including the deployment of up to 150,000 AI accelerators, including our GB 300, x AI and Humane also announced a partnership in which the two will jointly develop a network of world-class GPU data centers anchored by the flagship 500 megawatt facility. Blackwell gained further momentum in Q3. As GB 300 crossed over GB 200 and contributed roughly two-thirds of the total Blackwell revenue. The transition to GB 300 has been seamless. With production shipments to the majority to the major, cloud service providers, hyperscalers, and GP clouds and is already driving their growth. The Hopper platform in its thirteenth quarter since exception recorded approximately \$2 billion in revenue in Q3.

H '20 sales were approximately \$50 million. Sizable purchase orders never materialized in the quarter due to geopolitical issues and the increasingly competitive market in China. While we were disappointed in the current state, that prevents us from shipping more competitive data center compute products to China, we are committed to continued engagement with the US and China governments. And will continue to advocate for America's ability to compete around the world. To establish a sustainable leadership and position in AI computing, America must win. The support of every developer, and be the platform of choice for every commercial business including those in China. The Rubin platform is on track to ramp in the 2026.

Powered by seven chips, the Vera Rubin platform will once again deliver an x factor improvement in performance relative to Blackwell. We have received silicon back from our supply chain partners and are happy to report that NVIDIA Corporation teams across the world are executing the bring up beautifully. Rubin is our third generation rack scale system substantially redefined the manufacturability while remaining compatible with Grace Blackwell our supply chain data center ecosystem, and cloud partners have now mastered the build to installation process of NVIDIA Corporation's RAC architecture. Our ecosystem will be ready for a fast Rubin ramp. Our annual x factor performance lead increases performance per dollar while driving down computing cost for our customers.

The long useful life of NVIDIA Corporation's CUDA GPUs is a significant TCO advantage over accelerators. CUDA's compatibility and our massive installed base extend the life NVIDIA Corporation systems well beyond their original estimated useful life. For more than two decades, we have optimized the CUDA ecosystem, improving existing workloads, accelerating new ones, and increasing throughput with every software release. Most accelerators without CUDA and NVIDIA Corporation's time-tested and versatile app architecture became obsolete within a few years as model technologies evolve. Thanks to CUDA, the 100 GPUs we shipped six years ago are still running at full utilization today. Powered by vastly improved software stack.

We have evolved over the past twenty-five years from a gaming GPU company to now an AI data center infrastructure company. Our ability to innovate across the CPU, the GPU, networking, and software, and ultimately drive down cost per token is unmatched across the industry. Our networking business purpose built for AI, and now the largest in the world. Generated revenue of \$8.2 billion, up 162% year over year. With NVLink, InfiniBand, and Spectrum X Ethernet, all contributing to growth. We are winning in data center networking as the majority of AI deployments now include our switches with Ethernet GPU attach rates roughly on par with InfiniBand.

Meta, Microsoft, Oracle, and xAI are building gigawatt AI factories with Spectrum X Ethernet switches. And each will run its operating system of choice highlighting the flexibility and openness of our platform. We recently introduced SPECTUM Spectrum XGS, a scale across technology that enables gigascale AI factories. NVIDIA Corporation is the only company with AI scale up scale out, and scale across platforms, reinforcing our unique position in the market as the AI infrastructure provider. Customer interest in NVLink Fusion continues to grow. We announced a strategic collaboration with Suzuki in October where we will integrate Fuzitsu's CPUs and NVIDIA Corporation GPUs via NVLink Fusion. Connecting our large ecosystems.

We also announced a collaboration with Intel to develop multiple generations of custom data center and PC products connecting NVIDIA Corporation and Intel's ecosystems using NVLink. This week at supercomputing 25, Arm announced that it will be integrating NVLink IP for customers to build CPU SoCs that connect with NVIDIA Corporation. Currently on its fifth generation, NVLink is the only proven scale up technology available on the market today. In the latest MLPerf training results, Blackwell Ultra delivered 5x faster time to train than hopper. NVIDIA Corporation swept every benchmark. Notably, NVIDIA Corporation is the only training platform to ledge bridge f p four while meeting the MLPerf's strict accuracy standards.

In semianalysis, inference max benchmark, Blackwell achieved the highest performance and lowest total cost of ownership across every model and use case. Particularly important is Blackwell's NVLink's performance on a mixture of experts. The architecture for the world's most popular reasoning models. On DeepSeek, r one,

Blackwell delivered 10x higher performance per watt and 10x lower cost per token versus h 200.

A huge generational leap fueled by our extreme codesign approach NVIDIA Corporation Dynamo, an open source, low latency, modular inference framework has now been adopted by every major cloud service provider leveraging Dynamos enablement and disaggregated inference the resulting such as MOE models, increase in performance of complex AI models AWS, Google Cloud, Microsoft Azure, and OCI have boosted AI inference performance for enterprise cloud customers. We are working on a strategic partnership with OpenAI focused on helping them build and deploy at least 10 gigawatts of AI data centers. In addition, we have the opportunity to invest in the company. We serve OpenAI, through their cloud partners. Microsoft Azure, OCI, and CoreWeave.

We will continue to do so for the foreseeable future. As they continue to scale, we are delighted to support the company to add self build infrastructure, and we are working toward a definitive agreement and are excited to support OpenAI's growth. Yesterday, celebrated an announcement with Anthropic. For the first time, Anthropic is adopting NVIDIA Corporation and we are establishing a deep technology partnership to support Anthropic's fast growth. We will collaborate to optimize anthropic models for CUDA, and deliver the best possible performance, efficiency, and TCO. We will also optimize future NVIDIA Corporation architectures for anthropic workloads. Anthropic's compute commitment is initially including up to one gigawatt of compute capacity, with Grace Blackwell and Vera Rubin systems.

Our strategic investments in anthropic menstural, opening eye, reflection, thinking machines, and other represent partnerships. That grow the NVIDIA Corporation CUDA AI ecosystem and enable every model to run optimally on NVIDIA Corporation's everywhere. We will continue to invest strategically while preserving our disciplined approach to cash flow management. Physical AI is already a multibillion dollar business addressing a multitrillion dollar opportunity, and the next leg of growth for NVIDIA Corporation. Leading US manufacturers and robotics innovators are leveraging NVIDIA Corporation's three computer architecture to train on NVIDIA Corporation. Test on Omniverse computer, and deploy real world AI on Justin robotic computers.

PTC and Siemens introduced new services that bring Omniverse powered digital twin workflows to their extensive installed base of customers. Companies including Belden, Caterpillar, Foxconn, Lucid Motors, Toyota, TSMC, and Wistron are building Omniverse Digital Twin factories to accelerate AI driven manufacturing and automation. Agility robotics, Amazon robotics, Figure, and skilled at AI are building our platform, tapping offerings such as NVIDIA Corporation Cosmos World Foundation Models for development, Omniverse for simulation and validation, and Jetson two power next generation intelligent robots. We remain focused on building resiliency and

redundancy in our global supply chain. Last month, in partnership with TSMC, we celebrated the first Blackwell wafer produced on US soil.

We'll continue to work with Foxconn, Vistron, Amcor, Spill, and others to grow our presence in The US over the next four years. Gaming revenue was \$4.3 billion, up 30% year on year driven by strong demand as 42 million gamers, while thousands of fans packed the GeForce Gamer Festival in South Korea. To celebrate twenty-five years of GeForce. NVIDIA Corporation Pro Visualization has evolved into computers for engineers and developers. Whether for graphics, or for AI. Professional visualization revenue was \$760 million, up 56% year over year. Was another record. Growth was driven by DGX Spark. The world's smallest AI supercomputer. Built on a small configuration of Grace Blackwell.

Automotive revenue was \$592 million, up 32% year over year primarily driven by self-driving solutions. We are partnering with Uber to scale the world's largest level four ready autonomous fleet built on the new NVIDIA Corporation Hyperion I four robotaxi reference architecture. Moving to the rest of the p and l. GAAP gross margins were 73.4% and non GAAP gross margins was 73.6%. Exceeding our outlook. Gross margins increased sequentially due to our data center mix, improved cycle time, and cost structure. GAAP operating expenses were up 8% sequentially and up 11% on non GAAP basis. The growth was driven by infrastructure compute, as well as higher compensation and benefits in engineering development costs.

Non GAAP effective tax rate for the third quarter was just over 17%. Higher than our guidance of 16.5% due to the strong US revenue. On our balance sheet, inventory grew 32% quarter over quarter while supply commitments increased 63% sequentially. We are preparing for significant growth ahead and feel good about our ability to execute against our opportunity set. Okay. Let me turn to the outlook for the fourth quarter. Total revenue is expected to be \$65 billion plus or minus 2%. At the midpoint, our outlook implies 14% sequential growth driven by continued momentum in the Blackwell architecture. Consistent with last quarter, we are not assuming any data center compute revenue from China.

GAAP and non GAAP gross margins are expected to be 74.875% respectively. Plus or minus 50 basis points. Looking ahead, to fiscal year twenty-seven, input costs are on the rise but we are working to hold gross margins in the mid-seventies. GAAP and non GAAP operating expenses are expected to be approximately \$6.7 billion and \$5 billion respectively. GAAP and non GAAP other income and expenses are expected to be an income of approximately \$500 million, excluding gains and losses from non-marketable and publicly held equity securities. GAAP and non GAAP tax expense are expected to be 17%. Plus or minus 1% excluding any discrete items. At this time, let me turn the call over to Jensen.

For him to say a few words.

Jensen Huang: Thanks, Colette. There's been a lot of talk about an AI bubble. From our vantage point, we see something very different. As a reminder, NVIDIA Corporation is unlike any other accelerator. We excel at every phase of AI. From pre-training and post-training to inference. And with our two-decade in CUDA x acceleration libraries, we are also exceptional. At science and engineering simulations, computer graphics, structured data processing, to classical machine learning. The world is undergoing three massive platform shifts at once. The first time since the dawn of Moore's Law. NVIDIA Corporation is uniquely addressing each of the three transformations. The first transition is from CPU general-purpose computing to GPU accelerated computing. As Moore's Law slows.

The world has a massive investment in non-AI software. From data processing to science and engineering simulations. Representing hundreds of billions of dollars in cloud computing spend each year. Many of these applications, which ran once exclusively on CPUs, are now rapidly shifting to CUDA GPUs. Accelerated computing has reached a tipping point. Secondly, AI has also reached a tipping point. And is transforming existing applications while enabling entirely new ones. For existing applications, generative AI, is replacing classical machine learning in search ranking, recommender systems, ad targeting, click-through prediction, to content moderation, The very foundations of hyperscale infrastructure. Meta's gem a foundation model for ad recommendations trained on large-scale GPU clusters exemplifies this shift.

In Q2, Meta reported over a 5% increase in ad conversions on Instagram and 3% gain on Facebook feed. Driven by generative AI based JEM. Transitioning to generative AI. Represents substantial revenue gains for hyperscalers. Now a new wave is rising. Agentic AI systems. Capable of reasoning, planning, and using tools. From coding assistants like Cursor and QuadCode to radiology tools like iDoc, legal assistants like Harvey, and AI chauffeurs like Tesla FSD and Waymo, These systems mark the next frontier of computing. The fastest growing companies in the world today OpenAI, Anthropic, xAI, Google, Cursor, Lovable, Replit, Cognition AI, Open Evidence, a bridge Tesla, are pioneering agentic AI. So there are three massive platform shifts.

The transition to accelerated computing is foundational and necessary. Essential in a post-Moore's law era. The transition to generative AI is transformational, and necessary supercharging existing applications and business models. And the transition to agentic and physical AI will be revolutionary, giving rise to new applications, companies, products, services. As you consider infrastructure investments, consider these three fundamental dynamics. Each will contribute to infrastructure growth in the coming years. NVIDIA Corporation is chosen because our singular architecture enables all three transitions. And thus so for any form and modality of AI across all industries, across every phase of AI, across all of the diverse computing needs, in a cloud, and also from cloud to enterprise to robots. One architecture.

Toshiya, back to you.

Toshiya Hari: We will now open the call for questions. Operator, would you please poll for questions?

Sarah: Thank you. At this time, I would like to remind everyone, in order to ask a question, press star, then the number one on your telephone keypad. Thank you. Your first question comes from Joseph Moore with Morgan Stanley. Your line is open. Great. Thank you. I wonder if you could update us.

Joseph Moore: You talked about the \$500 billion of revenue for Blackwell plus Rubin. 'twenty five and 'twenty six at GTC. At that time, you talked about \$150 billion of that already having been shipped. So as the quarter's wrapped up, are those still kind of the general parameters that there's \$350 billion in the next kind of, you know, fourteen months or so. And, you know, I would assume over that time, you haven't seen all the demand that there is. There's possibility of upside to those numbers as we move forward?

Colette Kress: Yeah. Thanks, Joe. I'll start first with a response here on that. Yes. That's correct. We are working into our \$500 billion forecast. And we are on track for that as we have finished some of the quarters. And now we have several quarters now in front of us to take us through the end of calendar year '26. The number will grow. And we will achieve, I'm sure, additional needs for compute that will be shippable by fiscal year '26. So we shipped \$50 billion this quarter. But we would be not finished if we didn't say that we'll probably be taking more orders.

For example, just even today, our announcements with KSA and that agreement in itself is four to 600,000 more GPUs over three years. Anthropic is also not new. So there's definitely an opportunity for us to have more on top of the \$500 billion that we announced.

Sarah: The next question comes from C.J. Muse with Cantor Fitzgerald. Your line is open.

C.J. Muse: Yes. Good afternoon. Thank you for taking the question. There's clearly a great deal of consternation around the magnitude of AI infrastructure build outs and the ability to fund such plans in the ROI. Yet, you know, at the same time, you're talking about being sold out every stood up GP is taken. The AI world hasn't seen the enormous benefit yet know, from d 300, never mind Rubin. And Gemini three just announced Grok five coming soon.

And so the question is this, when you look at that as the backdrop, do you see a realistic path for supply to catch up with demand over the next twelve to eighteen months, or do you think it can extend beyond that time frame?

Jensen Huang: Well, as you know, we've done a really good job planning our supply chain. NVIDIA Corporation's supply chain basically includes every technology company in the world. And TSMC and their packaging and our memory vendors and memory partners and all of our system ODMs have done a really good job planning with us. And we were planning for a big year. You know, we've seen for some time the three transitions that I spoke about just a second ago, accelerated computing, from general-purpose computing and it's really important to recognize that AI is not just agentic AI, but generative AI is transforming the way that hyperscalers did the work that they used to do on CPUs.

Generative AI made it possible for them to move search and recommender systems and, you know, add recommendations and targeting. All of that has been generated has been moved to generative AI. And it's still transitioning. And so whether you install NVIDIA Corporation GPUs for data processing, or you did it for generative AI for your recommender system, or you're building it for agentic chatbots and the type of AIs that most people see when they think about AI, all of those applications are accelerated by NVIDIA Corporation. And so when you look at the totality of the spend, it's really important to think about each one of those layers. They're all growing. They're related, but not the same.

But the wonderful thing is that they all run on NVIDIA Corporation GPUs.

Simultaneously, because the quality of the AI models are improving so incredibly. The adoption of it in the different use cases, whether it's in code assistance, which NVIDIA Corporation uses fairly exhaustively, and we're not the only one. I mean, the fastest growing application in history combination of cursor and CliveCode and code OpenAI's codex and GitHub Copilot. These applications are the fastest growing in history. And it's not just used for software engineers. It's used by because of vibe coding, it's used by engineers and marketeers all over companies.

Supply chain planners, all over companies, And so I think that's just one example, the list goes on. You know, whether it's open evidence and work that they do in health care or the work that's being done in digital video editing runway. I mean, number of it really, really exciting start ups that are taking advantage of generative AI and agentic AI is growing quite rapidly, and not to mention all using it a lot more. And so all of these exponentials not to mention, you know, just today, I was reading a text from Dennis, and he was saying that pre-training and post-training are fully intact. You know?

And Gemini three takes advantage of the scaling laws, and got it received a huge jump in quality performance model performance. And so we're seeing all of these exponentials kind of running at the same time. And just always go back to first principles and think about what's happening from each one of the dynamics that I mentioned before. General-purpose computing to accelerated computing, generative AI replacing classical machine learning, and, of course, agentic AI, which is a brand

new category.

Sarah: The next question comes from Vivek Arya with Bank of America Securities. Your line is open.

Vivek Arya: Thanks for taking my question. I'm curious what are you making on NVIDIA Corporation content per gigawatt? In that \$500 billion number? Because we have heard, you know, numbers as low as \$25 billion per gigawatt of content to as high as \$30 or \$40 billion per gigawatt. So I'm curious what power and what dollar per gigawatt assumptions you are making as part of that \$500 billion number, And then longer term, Jensen, the three to \$4 trillion in data center by 2030 was mentioned. How much of that do you think will require vendor financing, and how much of that can be supported by cash flows of your large customers or governments or enterprises. Thank you.

Jensen Huang: In each generation, from Ampere to Hopper, from Hopper to Blackwell, Blackwell to Rubin, we are our a part of the data center increases. And hock regeneration was probably something along the lines of twenty some odd, twenty to twenty-five. I Blackwell generation, Grace Blackwell particularly, is probably 30 to 30, you know, say 30 plus or minus. And then Rubin is probably higher than that. And in each one of these generations, the speed up is x factors, And therefore their TCO, the customer TCO, improves by x factors, and the most important thing is in the end, you still only have one gigawatt of power.

You know, one gigawatt data center is one gig gigawatt power, and, therefore, performance per watt, the efficiency of your architecture is incredibly important. And the efficiency of your architecture can't be brute forced. There is no brute forcing about it. That one gigawatt translates directly. Your performance per watt translates directly absolutely directly to your revenues. Which is the reason why choosing the right architecture matters so much now. You know, the world doesn't have an excess of anything to squander. And so we have to be really, really you know, we use this concept called co-design.

Across our entire stack across the frameworks and models, across the entire data center, even power and cooling, optimized across the entire supply chain in our ecosystem. And so each generation our economic contribution will be greater Our value delivered will be greater. But the most important thing is our energy efficiency per watt is going to be extraordinary every single generation. With respect to growing into continuing to grow our customers financing is up to them. We are we see the opportunity to grow. For quite some time, and remember, today most of the focus has been on the hyperscalers.

And one of the areas that is really misunderstood about the hyperscalers is that the investment on NVIDIA Corporation GPUs not only improves their scale, speed, and

cost, for from general-purpose computing. That's number one, because Moore's Law has Moore's law scaling has really slowed. Moore's law is about driving cost down. It's about it's about deflationary cost, the incredible deflationary cost of computing over time. But that has slowed. Therefore, a new approach is necessary for them to keep driving the cost down. Going to NVIDIA Corporation GPU computing is really the best way to do so. The second is revenue boosting in their current business models. You know, recommender systems drive the world's hyperscalers.

Every single whether it's you know, watching short form videos or recommending books or recommending the next item in your basket to recommending ads to recommending news to rep it's all about recommenders. The world has the Internet has trillions of pieces of content How could they possibly figure out what to put in front of you in your little tiny screen unless they have really sophisticated recommender systems to do so. Well, that has gone generative AI. So the first two things that I just said hundreds of billions of dollars of CapEx is gonna have to be invested, is fully cash flow funded. What is above it, therefore, is AgenTik AI.

This is revenue is net new net new consumption, but it's also net new applications. And some of the applications I've mentioned before, but these are these new applications are also the fastest growing applications in history. Okay? So I think that I you're gonna see that once people start to appreciate what is actually happening under, you know, under the water, if you will, you know, from the simplistic view of what's happening to CapEx investment recognizing there's these three dynamics. And then lastly, remember we were just talking about the American CSPs. Each country will fund their own infrastructure. And you have multiple countries, You have multiple industries.

Most of the world's industries haven't really engaged AgenTic AI yet. And they're about to. You know? All the names of companies that you know we're working with know, whether it's autonomous vehicle companies or digital twins for physical AI for factories and the number of factories and warehouses being built around the world. Just the number of digital biology startups that are being funded so that we could accelerate drug discovery. All of those different industries are now getting engaged, they're gonna do their own fundraising. And so don't just look at the hyperscalers, as a way to build out for the future.

You gotta look at the world, you gotta look at all the different industries, and, you know, enterprise computing is gonna fund their own industry.

Sarah: The next question comes from Ben Reitzes with Melius. Your line is open.

Ben Reitzes: Hey, thanks a lot. Jensen, wanted to ask you about cash. Speaking of \$05 trillion you may generate about \$500 billion in free cash flow over the next couple of years. What are your plans for that cash? How much goes to buyback versus in the ecosystem? And how do you look at investing in the ecosystem? I think there's there's

just a lot of confusion out there about how these how these deals work and your criteria for doing those, like the Anthropic, the OpenAI's, etcetera. Thanks a lot.

Jensen Huang: Yeah. Appreciate the question. Of course, using cash to fund our growth No company has ever grown at the scale that we're talking about and have the connection and the depth and the breath of supply chain that NVIDIA Corporation has. The reason why our entire customer base can rely on us is because we've secured a really you know, really resilient supply chain and we have the balance sheet to support them. When we make purchases, our suppliers can take it to the bank. When we make when we make forecasts and we plan with them, they take us seriously. Because of our balance sheet. We're not we're not making up the offtake. We know what our offtake is.

And because they've been planning with us for so many years, our reputation and our credibility is incredible. And so it takes really strong balance sheet to do that. To support the level of growth and the rate of growth and the magnitude associated with that. So that's number one. The second thing, of course, we're gonna continue to do stock buyback. Buybacks. We're gonna continue to do that. But with respect to the investments, this is really, really important work that we do. All of the investments that we've done so far. Well all the week period is associated with expanding the reach of CUDA, expanding the ecosystem.

If you look at the work, investments that we did with OpenAI, of course, that relationship we've had since 2016. Delivered the first AI supercomputer ever made. To OpenAI. So we've had a close and wonderful relationship with OpenAI since then. And everything that OpenAI does runs on NVIDIA Corporation today. So all the clouds that they deploy in, whether it's training and inference, runs NVIDIA Corporation, and we love working with them. The partnership that we have with them is one so that we could work even deeper from a technical perspective so that we could support their accelerated growth This is a company that's growing incredibly fast. And don't just look at don't just look at no.

What is in the press. Look at all the ecosystem partners and all the developers that are connected to OpenAI. And they're all driving consumption of it. And the quality of the AI that's being produced huge step up since a year ago. So the quality of response is extraordinary. So we invest in OpenAI for a deep partnership and co-development to expand our ecosystem and to support their growth. And, of course, rather than giving up a share of our company, we get a share of their company.

And we invested in them in one of the most consequential once in a generation companies once in a generation company that we have a share And so I fully expect that investment to translate to extraordinary returns. Now in the case of Anthropic, this is the first time that Anthropic will be on NVIDIA Corporation's architecture. The first time NVIDIA Corporation will be Anthropic will be on NVIDIA Corporation's architecture

is the second most successful AI in the world, in terms of total number of users But in enterprise, they're doing incredibly well. ClotCode is doing incredibly well. Clot is doing incredibly well, all of the world's enterprise.

And now we have the opportunity to have a deep partnership with them and bringing Claude onto the NVIDIA Corporation platform. And so what do we have now? NVIDIA Corporation's architecture, taking a step back, NVIDIA Corporation's architecture NVIDIA Corporation's platform, is the singular platform in the world that runs every AI model. We run OpenAI, We run Anthropic. We run XAI. Because of our deep partnership with Elon and x AI, we were able to bring that opportunity to Saudi Arabia to the KSA so that humane could also be hosting opportunity for x AI. We run x AI. We run we run Gemini. We run thinking machines. Let's see. What else do we run? We run them all.

And so not to mention, we run the science models, the biology models, DNA models, gene models, chemical models, and all the different fields around the world. It's not just cognitive AI that the world uses. AI is impacting every single industry. And so we have the ability through the ecosystem investments that we make to partner with deeply partner on a technical basis with some of the best companies, most brilliant companies in the world, We are expanding the reach of our ecosystem and we're getting a share and investment in what will what will be a very successful company, oftentimes once in a generation company. And so that basic that's our that's our investment thesis.

Sarah: The next question comes from Jim Schneider with Goldman Sachs. Your line is open.

Jim Schneider: Afternoon. Thanks for taking my question. In the past, you've talked about roughly 40% of your shipments tied to AI inference I'm wondering as you look forward into next year, where do you expect that percentage could go in say a year's time? And can you maybe address the Rubin CPX product you expect to introduce next year and contextualize that? How big of the overall TAM you expect that can take and maybe talk about some of the target customer applications for that specific product? Thank you.

Jensen Huang: CPX is designed for long context type of workload generation. And so long context, basically, before you start generating answers, you have to read a lot. Basically, you know, long context. And it could be a bunch of PDFs, it could be watching a bunch of videos, studying three d images, so on and so forth. You have to you have to absorb the context. And so CPX is designed for long context type of workloads. And it's perf per dollars it's perf per dollar is excellent. It's perf for what is excellent. And which made me forget the first part of the question.

Colette Kress: Inprinting.

Jensen Huang: Oh, inference. Yeah. There are three scaling laws that are that are scaling at the same time. The first scaling law called pre-training, continues to be. Very effective. And the second is post-training. Post-training basically has found incredible algorithms for improving an AI's ability to break a problem down, and solve a problem step by step. And post-training is scaling exponentially. Basically, the more compute, you apply to a model, the smarter it is. The more intelligent it is. And then the third is inference. Inference because of chain of thought, because of reasoning capabilities, AIs are essentially reading, thinking, before it answers. And the amount of computation necessary as a result of those three things has gone completely exponential.

I think that it's hard to know exactly what the percentage will be at any given point in time and who. But, of course, our hope our hope is that inference is a very large part of the market. Because if inference is large, then what it suggests is that people are using it in more applications, and they're using it more frequently. And that's you know, we should all hope for inference to be very large. And this is where Grace Blackwell is just an order of magnitude better more advanced than anything in the world.

The second best platform is h 200, and it's very clear now that g b 300, g b 200, and g b 300, because of MP Link 72, the scale up network that we have, achieved. And you saw and Colette talked about in the seminar analysis benchmark it's the largest single inference benchmark ever done, And GB g b 200 m b link 72 is 10 times 10 to 15 times higher performance. And so that's a big step up. It's gonna take a long time before somebody is able to take that on. And our leadership there is surely multiyear. Yep. And so I think I'm hoping that inference becomes a very big deal.

Our leadership in inference is extraordinary.

Sarah: The next question comes from Timothy Arcuri with UBS. Your line is open.

Timothy Arcuri: Thanks a lot. Jensen, many of your customers are pursuing behind the meter power, but like what's the single biggest bottleneck that worries you that could constrain your growth? Is it power or maybe it's financing, or maybe it's, you know, something else like memory or even foundry? Thanks a lot.

Jensen Huang: Well, these are all issues and they're all constraints. And the reason for that, when you're growing at the rate that we are and the scale that we are, how could anything be easy? What NVIDIA Corporation is doing obviously has never been done before. And we've created a whole new industry. On the one hand, we are transitioning computing from general-purpose and classical or traditional computing accelerated computing and AI. That's on one hand. On the other hand, we created a whole new industry called AI factories. The idea that in order for software to run, you need these factories to generate it generate every single token instead of retrieving information that was pre created.

And so I think this whole transition requires extraordinary scale. And all the way from the supply chain of course, the supply chain, we have we have much better visibility and control over because obviously, we're incredibly good at managing our supply chain. We have great partners that we've worked with for thirty-three years. And so the supply chain part of it, we're quite confident. Now looking down our supply chain, we've now established partnerships with so many players in land and power and shell and, of course, financing. These things none of these things are easy. But they're all attractable, and they're all solvable things.

And the most important thing that we have to do is do a good job planning. We plan up the supply chain, down the supply chain, we've established a whole lot of partners and so we have a lot of routes to market. And very, you know, very importantly, our architecture has to deliver the best value to the customers that we have. And so at this point, you know, I'm I'm very confident that NVIDIA Corporation's architecture is the best performance per TCL. It is the best performance per watt and therefore, for any amount of energy that is delivered our architecture will drive the most revenues.

And I think the increasing rate of our success I think that we're more successful this year at this point than we were last year at this point. You know, the number of customers coming to us and the number of platforms coming to us after they've explored others is increasing, not decreasing. And so think the I think all of that is just, you know, all the things that I've been telling you over the years are really coming are coming true and or becoming evident.

Sarah: The next question comes from Stacy Rasgon with Bernstein Research. Your line is open.

Stacy Rasgon: Questions. Colette, I had some questions on margins. You said for next year, you're working to hold them in the mid-seventies. So I guess, first of all, what are the biggest cost increases? Is it just memory, or is it something else? What are you doing to work toward that? Is it how much is, like, you know, cost optimizations versus pre buys versus pricing. And then also, how should we think about OpEx growth next year given the revenues seem likely to grow materially? From where we're running right now?

Colette Kress: Stacy. Let me see if I can start with remembering where we were with the current fiscal year that we're in. Remember earlier this year, we indicated that through cost improvements and mix that we would exit the year in our gross margins in the mid-seventies. We achieved that. So now it's time for us to communicate where we are. And getting ready to also execute that in Q4. Working right now in terms of next year. Next year, there are input prices. That are well known in the industries that we need to work through. And our systems are by no means very easy to work with.

There are tremendous amount of components, many different parts of it as we think

about that. So we're taking all of that into account, but we do believe if we look at working again on cost improvements, cycle time, and mix, that we will work to try and hold at our gross margins in the mid-seventies. So that's our overall plan for gross margin. Your second question is around OpEx. And right now, our goal in terms of OpEx is to really make sure that we are innovating with our engineering teams, with all of our business teams, to create more and more systems for this market.

As you know, right now, have a new architecture coming out, and that means they are quite busy in order to meet that goal. And so we're gonna continue to see our investments on innovating more and more both our software, both our systems, and our hardware to do so. I'll leave it turn it to Jensen if he wants to add any couple more comments. Yeah.

Jensen Huang: I think that's spot on. I think the only thing that would add is remember that we plan, we forecast, we plan, and we negotiate with our supply chain well in advance. Our supply chain have known for quite a long time our requirements and they've known for quite a long time our demand, and we've been working with them and negotiating with them for quite a long time. And so I think the recent surge obviously quite significant, But remember, our supply chain has been working with us for a very long time.

It's a And so in many cases, we've secured a lot a lot of supply for ourselves, because, you know, obviously, they're working with the largest company in the world. In doing so. And we've also we've also been working closely with them on the financial aspects of it and securing forecasts and plans and so on and so forth. So I think all of that has worked out well for us.

Sarah: Your final question comes from the line of Aaron Rakers with Wells Fargo. Your line is open.

Aaron Rakers: Jensen, the question is for you. As you think about the entropic deal that was announced and just the overall breadth of your customers, I'm curious if your thoughts around the role that AI ASICs or dedicated play in these architecture build outs that has changed at all? Have you seen you know, I think you've been fairly adamant in the past that some of these some of these programs never really see deployments. But I'm I'm curious if we're at a point where maybe that's even changed more in favor of just GPU architecture. Thank you.

Jensen Huang: Yeah. Thank you very much. And I re I really appreciate the question. So first of all, you're competing against teams. You're you're excuse me, against a company. You're get competing against teams. And there are there just aren't that many teams in the world who are built who are extraordinary at building these incredibly complicated things. You know, back in the hopper day, and the ampere days, we would build one GPU. That's the definition of an accelerated AI system. But today,

we've gotta build entire racks, entire, you know, three different types of switches. A scale up, a scale out, and a scale across switch.

And it takes a lot more than one chip to build a compute node anymore. Everything about that computing system because AI needs to have memory, AI didn't used to have memory at all, Now it has to remember things. The amount of memory and context it has is gigantic. The memory the memory architecture implication is incredible. The diversity of models from mixture of experts to dense models to diffusion models to autoregressive, not to mention you know, biological models that obeys the laws of physics, the list of the list of different types of models have exploded in the last several years.

And so the challenge is the complexity of the problem is much higher, the diversity of AI models, is incredibly, incredibly large. And so this is where you know, if I will say the five things that makes us special, if you will. You know, the first thing I would say that makes us special is that we accelerate every phase of that transition. That's the first phase. That CUDA allows us to have CUDA x for transitioning from general-purpose of accelerated computing We are incredibly good at generative AI. We're incredibly good at agentic AI. So every single phase of that, every single layer of that transition, we are excellent at.

You can invest in one architecture, use it across the board, You can use what one architecture, and, not worry about the changes in the workload across those three phases. That's number one. Number two, we're excellent at every phase of AI. Everybody's always known that we're incredibly good pre-training. We're obviously very good at post-training. We're incredibly good as it turns out, at inference, because inference is really, really hard. How could thinking be easy? You know, people think that inference is one shot and therefore, it's easy. Anybody could approach the market that way. But it turns out to be the hardest of all because thinking, as it turns out, is quite hard.

We're great at every phase of AI, the second thing, The third thing is we're now the only architecture in the world that runs every AI model. Every frontier AI model we run open source AI models incredibly well. We run science models, biology models, robotics models, run every single model. We're the only architecture in the world that can claim that. It doesn't matter whether you're autoregressive or diffusion based. We run everything. And we run it for every major platform as I just mentioned. So we run every model. And then the fourth thing I would say is that in every cloud. The reason why developers love us is because we're literally everywhere. We're in every cloud.

We're in every we could even make you a little tiny cloud called DGX Spark. And so we're in every computer. We're everywhere from cloud to on-prem. To robotic systems. Edge devices, PCs, you name it. One architecture things just work. It's incredible. And then the last thing, this is probably the most important thing, the fifth thing, is if you are

a cloud service provider, if you're a new company like Humane, if you're a new company like CoreWeaver, Enscaler, Nebius, or OCI for that matter, The reason why NVIDIA Corporation is the best platform for you is because our off take is so diverse. We can help you with off take.

It's not about just putting a random ASIC into a data center. Where's the offtake coming from? Where's the diversity coming from? Where's the resilience coming from? The, you know, the versatility of the architecture coming from, the diversity of capability coming from. NVIDIA Corporation has such incredibly good offtake our ecosystem is so large. So these five things every phase of acceleration and transition, every phase of AI, every model, every cloud to on-prem, and, of course, finally, it all leads to offtake.

Sarah: Thank you. I will now turn the call to Toshiya Hari for closing remarks.

Toshiya Hari: In closing, please note we will be at the UBS Global and AI Conference on December 2. And our earnings call to discuss the results of our 2026 is scheduled for February 25. Thank you for joining us today. Operator, please go ahead and close the call. Thank you.

Sarah: This concludes today's conference call. May now disconnect.


This article is a transcript of this conference call produced for The Motley Fool. While we strive for our Foolish Best, there may be errors, omissions, or inaccuracies in this transcript. Parts of this article were created using Large Language Models (LLMs) based on The Motley Fool's insights and investing approach. It has been reviewed by our AI quality control systems. Since LLMs cannot (currently) own stocks, it has no positions in any of the stocks mentioned. As with all our articles, The Motley Fool does not assume any responsibility for your use of this content, and we strongly encourage you to do your own research, including listening to the call yourself and reading the company's SEC filings. Please see our Terms and Conditions (<https://www.fool.com/legal/terms-and-conditions/fool-rules>) for additional details, including our Obligatory Capitalized Disclaimers of Liability.

The Motley Fool has positions in and recommends Nvidia. The Motley Fool has a disclosure policy (<https://www.fool.com/legal/fool-disclosure-policy/>).

Advertisement

SPONSORED CONTENT

dianomi Advertise Here 



With E*TRADE, seamlessly trade across web, desktop, and mobile
E*TRADE from Morgan Stanley



5 Ways to Prepare for Unexpected Events
Charles Schwab



Firing Your Financial Advisor: The 5 Major Red Flags
Finance Advisors



Schwab's Market Outlook
Charles Schwab



Today's Markets: What Investors Should Know Now
Charles Schwab



How Savvy Investors Pay for Healthcare in Retirement [Free Guide]
Fisher Investments



Leverage TradeStation's platform to trade across asset classes.
TradeStation



5 Costly Mistakes to Avoid When Saving for College
Charles Schwab

Premium Investing Services

Invest better with The Motley Fool. Get stock recommendations, portfolio guidance, and more from The Motley Fool's premium services.

[View Premium Services \(/services/?ftm_cam=footer-services-sitewide-prospects\)](/services/?ftm_cam=footer-services-sitewide-prospects)